

## Load data and set up base autoencoder

```
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import fashion_mnist
from tensorflow.keras.layers import Input, Dense, Flatten, Reshape
from tensorflow.keras.models import Model

# Load Fashion MNIST data
(x_train, _), (x_test, _) = fashion_mnist.load_data()

# Normalize data to values between 0 and 1
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.

# Flatten the 28x28 images into vectors of size 784
x_train = x_train.reshape((x_train.shape[0], 28 * 28))
x_test = x_test.reshape((x_test.shape[0], 28 * 28))

# Define the dimensions
input_dim = 28 * 28 # 784 input neurons
encoding_dim = 64 # Compress to 64 features

# Define the input layer
input_layer = Input(shape=(input_dim,))

# Define the encoder
encoded = Dense(encoding_dim, activation='relu')(input_layer)

# Define the decoder
decoded = Dense(input_dim, activation='sigmoid')(encoded)

# Create the autoencoder model
autoencoder = Model(input_layer, decoded)

# Compile the model
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
```

Fit code using early stopping callback and run

-Stopped at epoch 91

```
from tensorflow.keras.callbacks import EarlyStopping

# Define EarlyStopping
early_stopping = EarlyStopping(monitor='val_loss',
                               patience=5, # Number of epochs with no improvement after which training will be stopped
                               restore_best_weights=True) # Restores model to best weights with the lowest validation loss

# Assuming x_train and x_test are your training and test datasets
autoencoder.fit(x_train, x_train, # For autoencoders, input and output are the same
                epochs=100, # Set a high number of epochs
                batch_size=256,
                shuffle=True,
                validation_data=(x_test, x_test),
                callbacks=[early_stopping]) # Add the early stopping callback


#Base code below

# Define the encoder model to get the compressed representation
encoder = Model(input_layer, encoded)

# Encode and decode some images from the test set
encoded_imgs = encoder.predict(x_test)
decoded_imgs = autoencoder.predict(x_test)

# Visualizing original and reconstructed images
n = 10 # Number of images to display
plt.figure(figsize=(20, 4))
for i in range(n):
    # Display original images
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
```

```
Epoch 90/100
235/235 ————— 1s 2ms/step - loss: 0.2687 - val_loss: 0.2718
Epoch 91/100
235/235 ————— 1s 2ms/step - loss: 0.2693 - val_loss: 0.2718
313/313 ————— 1s 1ms/step
313/313 ————— 1s 1ms/step
```



## TerminateOnNaN

```
[ ] from tensorflow.keras.callbacks import TerminateOnNaN

# Define the TerminateOnNaN callback
terminate_on_nan = TerminateOnNaN()

# Assuming x_train and x_test are your training and validation datasets
autoencoder.fit(x_train, x_train, # For autoencoders, input and output are the same
                epochs=100, # Set the number of epochs
                batch_size=256,
                shuffle=True,
                validation_data=(x_test, x_test),
                callbacks=[terminate_on_nan]) # Add the TerminateOnNaN callback

#Base Code Below


# Define the encoder model to get the compressed representation
encoder = Model(input_layer, encoded)

# Encode and decode some images from the test set
encoded_imgs = encoder.predict(x_test)
decoded_imgs = autoencoder.predict(x_test)


# Visualizing original and reconstructed images
n = 10 # Number of images to display
plt.figure(figsize=(20, 4))
for i in range(n):
    # Display original images
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # Display reconstructed images
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28), cmap='gray')
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```


Epoch 98/100

235/235  1s 3ms/step - loss: 0.2686 - val\_loss: 0.2717

Epoch 99/100








235/235  1s 3ms/step - loss: 0.2694 - val\_loss: 0.2716

Epoch 100/100

235/235  1s 3ms/step - loss: 0.2690 - val\_loss: 0.2716

313/313  0s 1ms/step

313/313  0s 1ms/step



## ModelCheckpoint

```
[ ] from tensorflow.keras.callbacks import ModelCheckpoint

# Define the ModelCheckpoint callback
checkpoint = ModelCheckpoint(filepath='autoencoder_best.keras', # File path to save the model
                             monitor='val_loss', # Metric to monitor
                             save_best_only=True, # Save only the best model (based on the monitored metric)
                             mode='min', # Minimize the monitored metric (e.g., validation loss)
                             save_weights_only=False, # Save the entire model (set to True to save only weights)
                             verbose=1) # Print a message when saving the model

# Assuming x_train and x_test are your training and validation datasets
autoencoder.fit(x_train, x_train, # For autoencoders, input and output are the same
               epochs=50, # Number of epochs
               batch_size=256,
               shuffle=True,
               validation_data=(x_test, x_test), # Validation data
               callbacks=[checkpoint]) # Add the ModelCheckpoint callback

#Base code below

# Define the encoder model to get the compressed representation
encoder = Model(input_layer, encoded)

# Encode and decode some images from the test set
encoded_imgs = encoder.predict(x_test)
decoded_imgs = autoencoder.predict(x_test)

# Visualizing original and reconstructed images
n = 10 # Number of images to display
plt.figure(figsize=(20, 4))
for i in range(n):
    # Display original images
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
```

219/235 ————— 0s 3ms/step - loss: 0.2681

Epoch 49: val\_loss did not improve from 0.27145

235/235 ————— 1s 4ms/step - loss: 0.2681 - val\_loss: 0.2715

Epoch 50/50









227/235 ————— 0s 2ms/step - loss: 0.2681

Epoch 50: val\_loss improved from 0.27145 to 0.27145, saving model to autoencoder\_best.keras

235/235 ————— 1s 2ms/step - loss: 0.2681 - val\_loss: 0.2714

313/313 ————— 0s 1ms/step

313/313 ————— 1s 2ms/step



## ReduceLROnPlateau

```
[ ] from tensorflow.keras.callbacks import ReduceLROnPlateau

# Define the ReduceLROnPlateau callback
reduce_lr = ReduceLROnPlateau(monitor='val_loss', # Metric to monitor
                              factor=0.5, # Factor by which the learning rate will be reduced (new_lr = lr * factor)
                              patience=3, # Number of epochs with no improvement after which learning rate will be reduced
                              min_lr=1e-6, # Lower bound for the learning rate
                              verbose=1) # Print message when the learning rate is reduced

# Assuming x_train and x_test are your training and validation datasets
autoencoder.fit(x_train, x_train, # For autoencoders, input and output are the same
               epochs=50, # Number of epochs
               batch_size=256,
               shuffle=True,
               validation_data=(x_test, x_test), # Validation data
               callbacks=[reduce_lr]) # Add the ReduceLROnPlateau callback

#Base code below


# Define the encoder model to get the compressed representation
encoder = Model(input_layer, encoded)

# Encode and decode some images from the test set
encoded_imgs = encoder.predict(x_test)
decoded_imgs = autoencoder.predict(x_test)

# Visualizing original and reconstructed images
n = 10 # Number of images to display
plt.figure(figsize=(20, 4))
for i in range(n):
    # Display original images
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # Display reconstructed images
```

235/235 ————— 1s 2ms/step - loss: 0.2681 - val\_loss: 0.2712 - learning\_rate: 1.0000e-06  
Epoch 48/50  
235/235 ————— 1s 2ms/step - loss: 0.2692 - val\_loss: 0.2712 - learning\_rate: 1.0000e-06  
Epoch 49/50  
235/235 ————— 1s 2ms/step - loss: 0.2688 - val\_loss: 0.2712 - learning\_rate: 1.0000e-06  
Epoch 50/50  
235/235 ————— 1s 2ms/step - loss: 0.2680 - val\_loss: 0.2712 - learning\_rate: 1.0000e-06  
313/313 ————— 0s 1ms/step  
313/313 ————— 0s 1ms/step





All four callbacks

```
[ ] from tensorflow.keras.callbacks import EarlyStopping
    from tensorflow.keras.callbacks import TerminateOnNaN
    from tensorflow.keras.callbacks import ModelCheckpoint
    from tensorflow.keras.callbacks import ReduceLRonPlateau

    # Define EarlyStopping
    early_stopping = EarlyStopping(monitor='val_loss',
                                   patience=10, # Number of epochs with no improvement after which training will be stopped
                                   restore_best_weights=True) # Restores model to best weights with the lowest validation loss

    # Define the TerminateOnNaN callback
    terminate_on_nan = TerminateOnNaN()


    # Define the ModelCheckpoint callback
    checkpoint = ModelCheckpoint(filepath='autoencoder_best.keras', # File path to save the model
                                  monitor='val_loss', # Metric to monitor
                                  save_best_only=True, # Save only the best model (based on the monitored metric)
                                  mode='min', # Minimize the monitored metric (e.g., validation loss)
                                  save_weights_only=False, # Save the entire model (set to True to save only weights)
                                  verbose=1) # Print a message when saving the model

    # Define the ReduceLRonPlateau callback
    reduce_lr = ReduceLRonPlateau(monitor='val_loss', # Metric to monitor
                                   factor=0.5, # Factor by which the learning rate will be reduced (new_lr = lr * factor)
                                   patience=3, # Number of epochs with no improvement after which learning rate will be reduced
                                   min_lr=1e-6, # Lower bound for the learning rate
                                   verbose=1) # Print message when the learning rate is reduced

    #Put all callbacks in
    autoencoder.fit(x_train, x_train, # For autoencoders, input and output are the same
                    epochs=100, # Number of epochs
                    batch_size=256,
                    shuffle=True,
                    validation_data=(x_test, x_test), # Validation data
                    callbacks=[reduce_lr, checkpoint, terminate_on_nan, early_stopping]) # Add the ReduceLRonPlateau callback

    #Base code below
```

```
Epoch 29: val_loss did not improve from 0.27118
235/235 ----- 1s 2ms/step - loss: 0.2689 - val_loss: 0.2712 - learning_rate: 1.0000e-06
Epoch 30/100
217/235 ----- 0s 2ms/step - loss: 0.2688
Epoch 30: val_loss did not improve from 0.27118
235/235 ----- 1s 3ms/step - loss: 0.2688 - val_loss: 0.2712 - learning_rate: 1.0000e-06
313/313 ----- 0s 1ms/step
313/313 ----- 0s 1ms/step
```



Reload data from `autoencoder_best.keras` and use on testing data.

I predict the results will be the same from the loaded model as the previous.

```
[ ] from tensorflow.keras.models import load_model

# Load the entire model
best_autoencoder = load_model('autoencoder_best.keras')

# Define the encoder model to get the compressed representation
encoder = Model(input_layer, encoded)

# Encode and decode some images from the test set
encoded_imgs = encoder.predict(x_test)
decoded_imgs = best_autoencoder.predict(x_test)

# Visualizing original and reconstructed images
n = 10 # Number of images to display
plt.figure(figsize=(20, 4))
for i in range(n):
    # Display original images
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # Display reconstructed images
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28), cmap='gray')
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```

```
ax = plt.subplot(2, n, i + 1 + n)
plt.imshow(decoded_imgs[i].reshape(28, 28), cmap='gray')
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()
```

313/313 ————— 0s 1ms/step  
313/313 ————— 1s 1ms/step

